

1. Business Understanding

This initial phase focuses on understanding the project objectives and requirements from a business perspective, and then converting this knowledge into a data mining problem definition, and a preliminary plan designed to achieve the objectives.

In this situation let's pretend we are a real estate agency in Boston MA and we are interested in purchasing some houses. We would like to know which houses are under value to help us narrow down the list and put in an accurate bid on a house.

Objective: Identify what makes a property valuable? What is a fair price for a house?

Dataset : Boston

Goal : Predict medv column in Test Dataset!

1. **crim** : per capita crime rate by town.
2. **zn** : proportion of residential land zoned for lots over 25,000 sq.ft.
3. **indus** : proportion of non-retail business acres per town.
4. **chas** : Charles River dummy variable (= 1 if tract bounds river; 0 otherwise).
5. **nox** : nitrogen oxides concentration (parts per 10 million).
6. **rm** : average number of rooms per dwelling.
7. **age** : proportion of owner-occupied units built prior to 1940.
8. **dis** : weighted mean of distances to five Boston employment centres.
9. **rad** : index of accessibility to radial highways.
10. **tax** : full-value property-tax rate per \$10,000.
11. **ptrato** : pupil-teacher ratio by town.
12. **black** : $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town.
13. **lstat** : lower status of the population (percent).
14. **medv** : median value of owner-occupied homes in \$1000s.

Load Library

```
In [1]: ▶ #import libraries for data handling
import os
import pandas as pd
import numpy as np

#import for visualization
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

#import for Linear regression
from sklearn.linear_model import LinearRegression
```

Load Data into Pandas Dataframe

```
In [2]: #Get Working Directory  
cwd = os.getcwd()  
cwd  
  
# Load the dataset  
file = cwd+'/train.xlsx'  
basetable1 = pd.read_excel(file)
```

```
In [3]: # peek preview into the data  
basetable1.head(6)
```

Out[3]:

	ID	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
3	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33
4	7	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.60	12.43
5	11	0.22489	12.5	7.87	0	0.524	6.377	94.3	6.3467	5	311	15.2	392.52	20.45

Here we see first 5 rows. Data is loaded Successfully!

2. Data Understanding (EDA)

Print a concise summary of a DataFrame.

This method prints information about a DataFrame including the index dtype and column dtypes, non-null values and memory usage.

```
In [4]: # Information on the Dataframe
print("\n\n", basetable1.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 333 entries, 0 to 332
Data columns (total 15 columns):
ID                333 non-null int64
crim              333 non-null float64
zn               333 non-null float64
indus            333 non-null float64
chas             333 non-null int64
nox              333 non-null float64
rm              333 non-null float64
age             333 non-null float64
dis             333 non-null float64
rad             333 non-null int64
tax             333 non-null int64
ptratio         333 non-null float64
black           333 non-null float64
lstat           333 non-null float64
medv           333 non-null float64
dtypes: float64(11), int64(4)
memory usage: 39.1 KB
```

None

```
In [5]: #how big is the data?
print(basetable1.size)
```

4995

Understand Correlation between Dependednt (medv) and other Features

```
In [6]: corr = basetable1.corr().tail(1)
corr.sort_values(by='medv',axis=1)
```

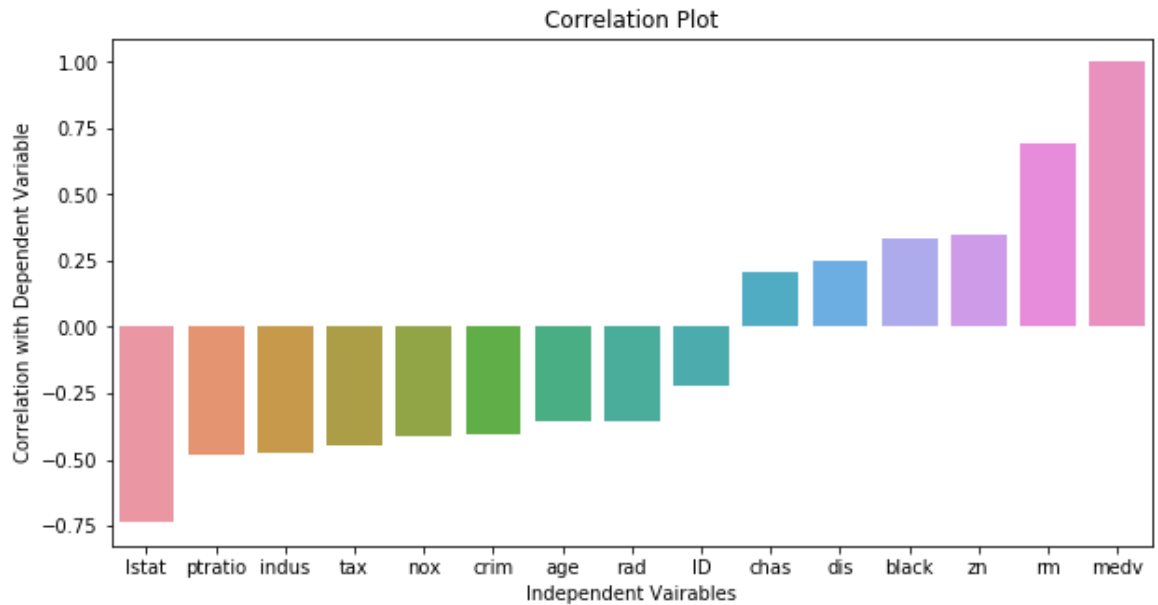
Out[6]:

	lstat	ptratio	indus	tax	nox	crim	age	rad
medv	-0.7386	-0.481376	-0.473932	-0.448078	-0.413054	-0.407454	-0.358888	-0.352251

```
In [7]: fig, ax = plt.subplots(figsize=(10,5))

plt.title("Correlation Plot")
plt.xlabel("Independent Vairables")
plt.ylabel("Correlation with Dependent Variable")

ax = sns.barplot(data=corr.sort_values(by='medv',axis=1))
```



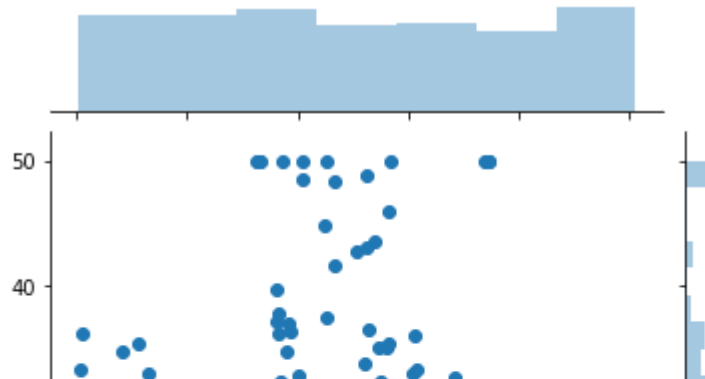
Correlation Based Inferences

1. ID Column though has Negative correlation but actually does not have any value to the Data. Eliminating at a later stage
2. rm : +vely MOST Impacting: average number of rooms per dwelling has High Correlation with medv (Dependent Variable)
3. lstat : -vely MOST Impacting: lower status of the population (percent) has High Correlation with medv (Dependent Variable)

```
In [8]: for index, columns in enumerate(basetable1.columns):
        svm = sns.jointplot(basetable1[basetable1.columns[index]], basetable1.medv,
```

C:\Users\kvibhaas\AppData\Local\Continuum\anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
        return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```




```
In [9]: # Set the palette to the "pastel" default palette:
sns.set_palette("pastel")
#Seaborn has six variations of its default color palette:: deep, muted, pastel
```

```
In [10]: from scipy import stats
#import stats for distribution graph
```


```
In [11]: sns.set(color_codes=True)
# set colour true
```

```
In [12]: x = np.random.gamma(6, size=200)
sns.distplot(x, kde=False, fit=stats.gamma);
```




```
In [13]: sns.distplot(basetable1.medv)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb7a2f10b8>
```




```
In [14]: sns.distplot(basetable1.medv, bins = 5, hist = True, rug = True)
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb7a3780b8>
```



```
In [15]: ➤ for index, columns in enumerate(basetable1.columns):
            sns.distplot(basetable1[basetable1.columns[index]])
            #svm = sns.jointplot(basetable1[basetable1.columns[index]], basetable1.medv
```



```
In [16]: ➤ basetable1.columns
```

```
Out[16]: Index(['ID', 'crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad',
                'tax', 'ptratio', 'black', 'lstat', 'medv'],
               dtype='object')
```

```
In [17]: # Python code to demonstrate the working of
# Log(a,Base)

import math
# Printing the Log base 5 of 14
print ("Logarithm base 10 of 14 is : ", end="")
print (math.log10(14))

#Length of train dataset
len(basetable1)

#Log of Length
print (math.log10(len(basetable1)))


# Struge formula
print(1 + 3.322*(math.log10(len(basetable1))))

# round off the bin size
bin_size = round((1 + 3.322*(math.log10(len(basetable1)))))

Logarithm base 10 of 14 is : 1.146128035678238
2.5224442335063197
9.379559743707995
```

```
In [18]: sns.distplot(basetable1.medv, bins = bin_size, hist = True, rug = True)
# number of bins are calculated as per Struge's rule  $K = 1 + 3.322 \log N$ 
```

Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb78bc2668>



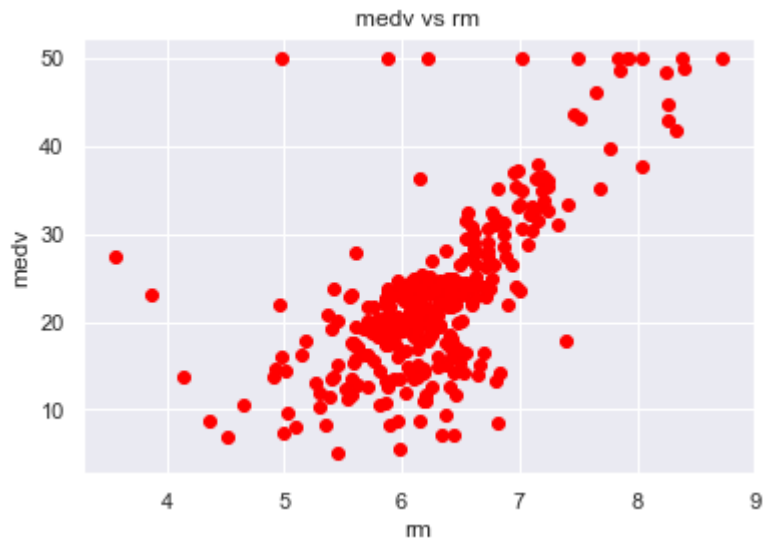
Linear Regression Assumptions

1. Linear relationship between target and features
2. No outliers
3. No high-leverage points
4. Homoscedasticity of error terms
5. Uncorrelated error terms
6. Independent features

#1 Linear Relationship Between Target & Features

```
In [19]: x = basetable1.rm  
y = basetable1.medv
```

```
In [20]: plt.scatter(x,y,color = 'red')  
#plot.plot(xTrain, linearRegressor.predict(xTrain), color = 'blue')  
plt.title('medv vs rm')  
plt.xlabel('rm')  
plt.ylabel('medv')  
plt.show()
```




```
In [21]: corr = basetable1.corr()
print(corr)
```

	ID	crim	zn	indus	chas	nox	
rm \							
ID	1.000000	0.456312	-0.155639	0.421978	0.007958	0.440185	-0.1127
90							
crim	0.456312	1.000000	-0.210913	0.422228	-0.041195	0.463001	-0.3101
80							
zn	-0.155639	-0.210913	1.000000	-0.518679	-0.024442	-0.501990	0.3281
97							
indus	0.421978	0.422228	-0.518679	1.000000	0.037496	0.750087	-0.4403
65							
chas	0.007958	-0.041195	-0.024442	0.037496	1.000000	0.080275	0.1122
51							
nox	0.440185	0.463001	-0.501990	0.750087	0.080275	1.000000	-0.3385
15							
rm	-0.112790	-0.310180	0.328197	-0.440365	0.112251	-0.338515	1.0000
00							
age	0.257300	0.379034	-0.544513	0.638378	0.068286	0.736000	-0.2485
73							
dis	-0.356461	-0.397067	0.637142	-0.702327	-0.081834	-0.769364	0.2691
91							
rad	0.707526	0.666636	-0.303663	0.569779	0.007714	0.612180	-0.2727
83							
tax	0.686246	0.617081	-0.311180	0.708313	-0.021826	0.670722	-0.3569
87							
ptratio	0.309838	0.313409	-0.380449	0.391087	-0.125067	0.192513	-0.3669
27							
black	-0.271619	-0.475796	0.168130	-0.335049	0.062029	-0.369416	0.1552
02							
lstat	0.281953	0.532077	-0.388112	0.614155	-0.050055	0.598874	-0.6157
47							
medv	-0.221694	-0.407454	0.344842	-0.473932	0.204390	-0.413054	0.6895
98							
	age	dis	rad	tax	ptratio	black	lst
at \							
ID	0.257300	-0.356461	0.707526	0.686246	0.309838	-0.271619	0.2819
53							
crim	0.379034	-0.397067	0.666636	0.617081	0.313409	-0.475796	0.5320
77							
zn	-0.544513	0.637142	-0.303663	-0.311180	-0.380449	0.168130	-0.3881
12							
indus	0.638378	-0.702327	0.569779	0.708313	0.391087	-0.335049	0.6141
55							
chas	0.068286	-0.081834	0.007714	-0.021826	-0.125067	0.062029	-0.0500
55							
nox	0.736000	-0.769364	0.612180	0.670722	0.192513	-0.369416	0.5988
74							
rm	-0.248573	0.269191	-0.272783	-0.356987	-0.366927	0.155202	-0.6157
47							
age	1.000000	-0.764208	0.447380	0.511893	0.259293	-0.268054	0.5888
34							
dis	-0.764208	1.000000	-0.477610	-0.529539	-0.231101	0.284374	-0.5059
39							
rad	0.447380	-0.477610	1.000000	0.903562	0.470849	-0.406405	0.4845

68								
tax	0.511893	-0.529539	0.903562	1.000000	0.467437	-0.406477	0.5444	
85								
ptratio	0.259293	-0.231101	0.470849	0.467437	1.000000	-0.164614	0.3748	
02								
black	-0.268054	0.284374	-0.406405	-0.406477	-0.164614	1.000000	-0.3566	
93								
lstat	0.588834	-0.505939	0.484568	0.544485	0.374802	-0.356693	1.0000	
00								
medv	-0.358888	0.249422	-0.352251	-0.448078	-0.481376	0.336660	-0.7386	
00								

	medv
ID	-0.221694
crim	-0.407454
zn	0.344842
indus	-0.473932
chas	0.204390
nox	-0.413054
rm	0.689598
age	-0.358888
dis	0.249422
rad	-0.352251
tax	-0.448078
ptratio	-0.481376
black	0.336660
lstat	-0.738600
medv	1.000000

```
In [23]: fig = plt.figure(figsize=(20,10))
ax = fig.add_subplot(111)
cax = ax.matshow(corr,cmap='coolwarm', vmin=-1, vmax=1)
fig.colorbar(cax)
ticks = np.arange(0,len(basetable1.columns),1)
ax.set_xticks(ticks)
plt.xticks(rotation=90)
ax.set_yticks(ticks)
ax.set_xticklabels(basetable1.columns)
ax.set_yticklabels(basetable1.columns)
plt.show()
```



3. Data Preparation

In [24]: `basetable1.head()`

Out[24]:

	ID	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
3	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33
4	7	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.60	12.43

```
In [25]: # Drop ID
basetable2 = basetable1.drop('ID',axis=1)

print("ID Column Dropped from Dataframe")

basetable2.head()
```

ID Column Dropped from Dataframe

Out[25]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	m
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	2
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	2
2	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	3
3	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	3
4	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.60	12.43	2

Split Data into Test & Train

Benefit to splitting a dataset into some ratio of training and testing subsets for a learning algorithm

- **Motivation:** we need a way to choose between machine learning models and our goal is to estimate likely performance of a model on out-of-sample data.
- **Initial idea:** we can train and test on the same data. However this will cause overfitting. As the number of features in a dataset increases the problem will increase
- **Alternative idea:** we can use train/test split. We can split the dataset into two pieces so that the model can be trained and tested on different data. Then, testing accuracy is a better estimate than training accuracy of out-of-sample performance.

```
In [26]: ▶ predictor = basetable2.drop('medv', axis = 1)
print(" Dependent variable : 'medv' Column removed from features")
predictor.head()
```

Dependent variable : 'medv' Column removed from features

Out[26]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
3	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33
4	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.60	12.43

```
In [30]: ▶ target = basetable2['medv']
print(" Target variable : 'medv' Column retained from features")
target.head()
```

Target variable : 'medv' Column retained from features

Out[30]:

0	24.0
1	21.6
2	33.4
3	36.2
4	22.9

Name: medv, dtype: float64

train_test_split

```
In [31]: ▶ from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(predictor, target, test_s

# Success
print("Training and testing split by 70/30 was successful")
```

Training and testing split by 70/30 was successful

```
In [32]: ▶ print("Training Predictor dimension :",x_train.shape)
print("Training Target dimension :",y_train.shape)
print("Test Predictor dimension :",x_test.shape)
print("Test Target dimension :",y_test.shape)
```

```
Training Predictor dimension : (233, 13)
Training Target dimension : (233,)
Test Predictor dimension : (100, 13)
Test Target dimension : (100,)
```

4. Model

lm1 : Raw data only removing ID

```
In [33]: ▶ from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

#Model Training
lm1 = LinearRegression(fit_intercept=True,normalize=False)
print("Parameters of Linear Regressor function : ",lm1.get_params)

#Model Training
lm1.fit(x_train,y_train)
# model is trained on raw data
```

```
Parameters of Linear Regressor function : <bound method BaseEstimator.get_
params of LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)>
```

```
Out[33]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

Prediction of Y based on test sample

```
In [34]: ▶ y_pred = lm1.predict(x_test)
print("Total number of predicted values = ",y_pred.shape)
```

```
Total number of predicted values = (100,)
```

```
In [35]: ▶ # The coefficients
print(lm1.coef_)
```


```
[-8.58952519e-02  5.98720246e-02 -1.28016377e-02  4.13703504e+00
 -1.39288018e+01  3.09043992e+00 -6.73962111e-03 -1.83154926e+00
  3.75417497e-01 -1.38415373e-02 -8.73159687e-01  1.46218453e-02
 -6.10464808e-01]
```

```
In [ ]: ▶
```

```
In [48]: # The coefficients
plt.plot( lm1.coef_)

plt.title("Coefficient Plot")
plt.xlabel("Predictors")
plt.ylabel("Coefficient")
#crim  zn  indus  chas  nox  rm  age  dis  rad  tax  ptratio  black  lstat
```

Out[48]: Text(0, 0.5, 'Coefficient')



5. Model Evaluation

1. root-mean-square error (RMSE) for the Model

2. R-Squared for the Model

```
In [50]: from math import sqrt

#Calculate root-mean-square error (RMSE):
print("R-Squared for the above model : ",r2_score(y_test,y_pred)*100,"%")

#Calculate R-squared for the Model:
print("\nroot-mean-square error (RMSE) for the model is : ",sqrt(mean_squared_error(y_test,y_pred)))

R-Squared for the above model : 78.08137322900414 %

root-mean-square error (RMSE) for the model is : 4.362458566264191
```

1. Using Statmodels.api to train the model

2 Print Summary for the Model


```
In [51]: import statsmodels.api as sm

model = sm.OLS(y_train,x_train).fit()
model.summary()
```

Out[51]:

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.954
Model:	OLS	Adj. R-squared:	0.951
Method:	Least Squares	F-statistic:	348.0
Date:	Tue, 17 Dec 2019	Prob (F-statistic):	1.05e-138
Time:	07:45:33	Log-Likelihood:	-714.65
No. Observations:	233	AIC:	1455.
Df Residuals:	220	BIC:	1500.
Df Model:	13		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
crim	-0.0408	0.068	-0.597	0.551	-0.175	0.094
zn	0.0612	0.023	2.710	0.007	0.017	0.106
indus	-0.0192	0.095	-0.203	0.840	-0.206	0.167
chas	4.5415	1.631	2.784	0.006	1.326	7.757
nox	1.7283	5.583	0.310	0.757	-9.276	12.732
rm	5.2955	0.504	10.509	0.000	4.302	6.289
age	-0.0180	0.022	-0.821	0.412	-0.061	0.025
dis	-1.2825	0.343	-3.744	0.000	-1.958	-0.607
rad	0.2140	0.106	2.015	0.045	0.005	0.423
tax	-0.0096	0.006	-1.592	0.113	-0.022	0.002
ptratio	-0.2779	0.179	-1.556	0.121	-0.630	0.074
black	0.0214	0.004	5.036	0.000	0.013	0.030
lstat	-0.5250	0.079	-6.679	0.000	-0.680	-0.370

Omnibus:	100.684	Durbin-Watson:	1.930
Prob(Omnibus):	0.000	Jarque-Bera (JB):	455.674
Skew:	1.707	Prob(JB):	1.13e-99
Kurtosis:	8.939	Cond. No.	9.01e+03

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.01e+03. This might indicate that there are strong multicollinearity or other numerical problems.

coef std err t P>|t| [0.025 0.975] crim -0.0408 0.068 -0.597 0.551 -0.175 0.094 indus -0.0192 0.095
-0.203 0.840 -0.206 0.167 nox 1.7283 5.583 0.310 0.757 -9.276 12.732 age -0.0180 0.022 -0.821
0.412 -0.061 0.025 tax -0.0096 0.006 -1.592 0.113 -0.022 0.002

zn 0.0612 0.023 2.710 0.007 0.017 0.106

chas 4.5415 1.631 2.784 0.006 1.326 7.757

rm 5.2955 0.504 10.509 0.000 4.302 6.289

dis -1.2825 0.343 -3.744 0.000 -1.958 -0.607 rad 0.2140 0.106 2.015 0.045 0.005 0.423

ptratio -0.2779 0.179 -1.556 0.121 -0.630 0.074 black 0.0214 0.004 5.036 0.000 0.013 0.030 lstat
-0.5250 0.079 -6.679 0.000 -0.680 -0.370

rad : index of accessibility to radial highways. tax : full-value property-tax rate per \$10,000.

High multi colinearity

```
In [53]: # Create New Feature Tax_Rad  
basetable2['tax_rad'] = basetable2.tax * basetable1.rad  
basetable2.head()
```

Out[53]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	m
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	2
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	2
2	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	3
3	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	3
4	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.60	12.43	2

```
In [54]: #drop Tax and Rad as they are highly multicollinear
basetable3 = basetable2.drop(['tax', 'rad'],axis=1)

print("tax & rad Column Dropped from Dataframe")

basetable3.head()
```

tax & rad Column Dropped from Dataframe

Out[54]:

	crim	zn	indus	chas	nox	rm	age	dis	ptratio	black	lstat	medv	Tax_F
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	15.3	396.90	4.98	24.0	2
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	17.8	396.90	9.14	21.6	4
2	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	18.7	394.63	2.94	33.4	6
3	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	18.7	396.90	5.33	36.2	6
4	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	15.2	395.60	12.43	22.9	15

```
In [55]: # Drop indus, nox, crim, age Based on High p Value

basetable3 = basetable2.drop(['indus', 'nox', 'crim', 'age'],axis=1)

print("indus, nox, crim, age Column Dropped from Dataframe")

basetable3.head()
```

indus, nox, crim, age Column Dropped from Dataframe

Out[55]:

	zn	chas	rm	dis	rad	tax	ptratio	black	lstat	medv	Tax_Rad
0	18.0	0	6.575	4.0900	1	296	15.3	396.90	4.98	24.0	296
1	0.0	0	6.421	4.9671	2	242	17.8	396.90	9.14	21.6	484
2	0.0	0	6.998	6.0622	3	222	18.7	394.63	2.94	33.4	666
3	0.0	0	7.147	6.0622	3	222	18.7	396.90	5.33	36.2	666
4	12.5	0	6.012	5.5605	5	311	15.2	395.60	12.43	22.9	1555

```
In [58]: basetable3 = basetable3.rename(columns={'Tax_Rad': 'tax_rad'})
```

```
In [59]: basetable3.head()
```

Out[59]:

	zn	chas	rm	dis	rad	tax	ptratio	black	lstat	medv	tax_rad
0	18.0	0	6.575	4.0900	1	296	15.3	396.90	4.98	24.0	296
1	0.0	0	6.421	4.9671	2	242	17.8	396.90	9.14	21.6	484
2	0.0	0	6.998	6.0622	3	222	18.7	394.63	2.94	33.4	666
3	0.0	0	7.147	6.0622	3	222	18.7	396.90	5.33	36.2	666
4	12.5	0	6.012	5.5605	5	311	15.2	395.60	12.43	22.9	1555

In []: ▶

Model Attempt 2

```
In [60]: ▶ #get Predictor Dataframe
predictor = basetable3.drop('medv', axis = 1)
print(" Dependent variable : 'medv' Column removed from features")
predictor.head()

#get Target Dataframe
target = basetable3['medv']
target.head()

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(predictor, target, test_s

# Success
print("Training and testing split by 70/30 was successful")

print("Training Predictor dimension :",x_train.shape)
print("Training Target dimension :",y_train.shape)
print("Test Predictor dimension :",x_test.shape)
print("Test Target dimension :",y_test.shape)

Dependent variable : 'medv' Column removed from features
Training and testing split by 70/30 was successful
Training Predictor dimension : (233, 10)
Training Target dimension : (233,)
Test Predictor dimension : (100, 10)
Test Target dimension : (100,)
```

```
In [61]: #import library
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

#Model Training
lm2 = LinearRegression(fit_intercept=True,normalize=False)
print("Parameters of Linear Regressor function : ",lm2.get_params())

#Model Training
lm2.fit(x_train,y_train)


#Predict
y_pred = lm2.predict(x_test)
print("Total number of predicted values = ",y_pred.shape)

# The coefficients
plt.plot(lm2.coef_)
```

Parameters of Linear Regressor function : <bound method BaseEstimator.get_params of LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)>

Total number of predicted values = (100,)

Out[61]: [<matplotlib.lines.Line2D at 0x2bb7742d358>]



0zn 1chas 2rm 3dis 4rad 5tax 6ptratio 7black 8lstat 9tax_rad chas : Charles River dummy variable (= 1 if tract bounds river; 0 otherwise). rm : average number of rooms per dwelling.

Room and Charles river has highest coefficient

Model Evaluation

In [62]: ▶ *#Model Evaluation*

```
from math import sqrt
```

```
#Calculate root-mean-square error (RMSE):
```

```
print("R-Squared for the above model : ",r2_score(y_test,y_pred)*100,"%")
```

```
#Calculate R-squared for the Model:
```

```
print("\nroot-mean-square error (RMSE) for the model is : ",sqrt(mean_squared
```

R-Squared for the above model : 59.77538914536744 %

root-mean-square error (RMSE) for the model is : 5.894945770741745

```
In [63]: ▶ model = sm.OLS(y_train,x_train).fit()  
model.summary()
```

Out[63]: OLS Regression Results

Dep. Variable:	medv	R-squared:	0.964
Model:	OLS	Adj. R-squared:	0.962
Method:	Least Squares	F-statistic:	596.7
Date:	Tue, 17 Dec 2019	Prob (F-statistic):	7.03e-155
Time:	08:26:03	Log-Likelihood:	-683.65
No. Observations:	233	AIC:	1387.
Df Residuals:	223	BIC:	1422.
Df Model:	10		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
zn	0.0776	0.019	4.148	0.000	0.041	0.114
chas	3.0757	1.495	2.058	0.041	0.130	6.021
rm	4.8211	0.379	12.718	0.000	4.074	5.568
dis	-1.1184	0.238	-4.695	0.000	-1.588	-0.649
rad	1.1862	0.380	3.119	0.002	0.437	1.935
tax	-0.0029	0.005	-0.600	0.549	-0.013	0.007
ptratio	-0.3249	0.154	-2.111	0.036	-0.628	-0.022
black	0.0158	0.004	4.214	0.000	0.008	0.023
lstat	-0.4835	0.060	-8.089	0.000	-0.601	-0.366
tax_rad	-0.0016	0.001	-2.829	0.005	-0.003	-0.000

Omnibus:	77.703	Durbin-Watson:	2.097
Prob(Omnibus):	0.000	Jarque-Bera (JB):	303.603
Skew:	1.317	Prob(JB):	1.18e-66
Kurtosis:	7.933	Cond. No.	4.24e+04

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.24e+04. This might indicate that there are strong multicollinearity or other numerical problems.


```
In [65]: ► basetable4 = basetable3.drop(['tax', 'rad'],axis=1)
basetable4.head()
```

Out[65]:

	zn	chas	rm	dis	ptratio	black	lstat	medv	tax_rad
0	18.0	0	6.575	4.0900	15.3	396.90	4.98	24.0	296
1	0.0	0	6.421	4.9671	17.8	396.90	9.14	21.6	484
2	0.0	0	6.998	6.0622	18.7	394.63	2.94	33.4	666
3	0.0	0	7.147	6.0622	18.7	396.90	5.33	36.2	666
4	12.5	0	6.012	5.5605	15.2	395.60	12.43	22.9	1555

Model Attempt 3

```
In [66]: ► #get Predictor Dataframe
predictor = basetable4.drop('medv', axis = 1)
print(" Dependent variable : 'medv' Column removed from features")
predictor.head()

#get Target Dataframe
target = basetable4['medv']
target.head()

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(predictor, target, test_s

# Success
print("Training and testing split by 70/30 was successful")

print("Training Predictor dimension :",x_train.shape)
print("Training Target dimension :",y_train.shape)
print("Test Predictor dimension :",x_test.shape)
print("Test Target dimension :",y_test.shape)
```

```
Dependent variable : 'medv' Column removed from features
Training and testing split by 70/30 was successful
Training Predictor dimension : (233, 8)
Training Target dimension : (233,)
Test Predictor dimension : (100, 8)
Test Target dimension : (100,)
```

```

In [67]: #import library
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

#Model Training
lm3 = LinearRegression(fit_intercept=True,normalize=False)
print("Parameters of Linear Regressor function : ",lm3.get_params)

#Model Training
lm3.fit(x_train,y_train)

#Predict
y_pred = lm3.predict(x_test)
print("Total number of predicted values = ",y_pred.shape)


# The coefficients
plt.plot(lm3.coef_)

```

Parameters of Linear Regressor function : <bound method BaseEstimator.get_params of LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)>

Total number of predicted values = (100,)

Out[67]: [<matplotlib.lines.Line2D at 0x2bb78d7a080>]



Ozn 1chas 2rm 3dis 4ptratio 5black 6lstat 7tax_rad

Model Evaluation

In [68]: ▶ *#Model Evaluation*

```
from math import sqrt
```

```
#Calculate root-mean-square error (RMSE):
```

```
print("R-Squared for the above model : ",r2_score(y_test,y_pred)*100,"%")
```

```
#Calculate R-squared for the Model:
```

```
print("\nroot-mean-square error (RMSE) for the model is : ",sqrt(mean_squared
```

```
R-Squared for the above model : 70.86083121002729 %
```

```
root-mean-square error (RMSE) for the model is : 5.205159161462403
```

R squared is still low 70% and RMSE is high 5.2

```
In [69]: > model = sm.OLS(y_train,x_train).fit()  
model.summary()
```

Out[69]: OLS Regression Results

Dep. Variable:	medv	R-squared:	0.959
Model:	OLS	Adj. R-squared:	0.958
Method:	Least Squares	F-statistic:	660.7
Date:	Tue, 17 Dec 2019	Prob (F-statistic):	1.20e-151
Time:	08:33:15	Log-Likelihood:	-701.85
No. Observations:	233	AIC:	1420.
Df Residuals:	225	BIC:	1447.
Df Model:	8		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
zn	0.0352	0.020	1.764	0.079	-0.004	0.075
chas	2.9610	1.290	2.296	0.023	0.420	5.502
rm	5.8555	0.379	15.468	0.000	5.110	6.601
dis	-0.8757	0.257	-3.401	0.001	-1.383	-0.368
ptratio	-0.5564	0.156	-3.574	0.000	-0.863	-0.250
black	0.0138	0.004	3.207	0.002	0.005	0.022
lstat	-0.4786	0.059	-8.121	0.000	-0.595	-0.362
tax_rad	-3.185e-05	6.79e-05	-0.469	0.640	-0.000	0.000

Omnibus:	94.266	Durbin-Watson:	1.895
Prob(Omnibus):	0.000	Jarque-Bera (JB):	446.643
Skew:	1.559	Prob(JB):	1.03e-97
Kurtosis:	9.024	Cond. No.	3.18e+04

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.18e+04. This might indicate that there are strong multicollinearity or other numerical problems.

tax_rad is not helping us. this is in significant. lets remove this and create model again. also Chas is zn has p value

Model Attempt 4

```
In [71]: ► basetable5 = basetable4.drop(['tax_rad', 'zn'],axis=1)
basetable5.head()
```

Out[71]:

	chas	rm	dis	ptratio	black	lstat	medv
0	0	6.575	4.0900	15.3	396.90	4.98	24.0
1	0	6.421	4.9671	17.8	396.90	9.14	21.6
2	0	6.998	6.0622	18.7	394.63	2.94	33.4
3	0	7.147	6.0622	18.7	396.90	5.33	36.2
4	0	6.012	5.5605	15.2	395.60	12.43	22.9

```
In [72]: ► #get Predictor Dataframe
predictor = basetable5.drop('medv', axis = 1)
print(" Dependent variable : 'medv' Column removed from features")
predictor.head()

#get Target Dataframe
target = basetable5['medv']
target.head()

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(predictor, target, test_s

# Success
print("Training and testing split by 70/30 was successful")

print("Training Predictor dimension :",x_train.shape)
print("Training Target dimension :",y_train.shape)
print("Test Predictor dimension :",x_test.shape)
print("Test Target dimension :",y_test.shape)
```

```
Dependent variable : 'medv' Column removed from features
Training and testing split by 70/30 was successful
Training Predictor dimension : (233, 6)
Training Target dimension : (233,)
Test Predictor dimension : (100, 6)
Test Target dimension : (100,)
```

```
In [73]: #import library
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

#Model Training
lm4 = LinearRegression(fit_intercept=True,normalize=False)
print("Parameters of Linear Regressor function : ",lm4.get_params)

#Model Training
lm4.fit(x_train,y_train)


#Predict
y_pred = lm4.predict(x_test)
print("Total number of predicted values = ",y_pred.shape)

# The coefficients
plt.plot(lm4.coef_)
```

Parameters of Linear Regressor function : <bound method BaseEstimator.get_params of LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)>

Total number of predicted values = (100,)

Out[73]: [<matplotlib.lines.Line2D at 0x2bb7795ca58>]



Model Evaluation

```
In [102]: ▶ #Model Evaluation

from math import sqrt

#Calculate root-mean-square error (RMSE):
print("R-Squared for the above model : ",r2_score(y_test,y_pred)*100,"%")

#Calculate R-squared for the Model:
print("\nroot-mean-square error (RMSE) for the model is : ",sqrt(mean_squared

R-Squared for the above model :  -84.6901264244536 %

root-mean-square error (RMSE) for the model is :  11.520263468475049
```

```
In [75]: ▶ model = sm.OLS(y_train,x_train).fit()
model.summary()
```

Out[75]:

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.962
Model:	OLS	Adj. R-squared:	0.961
Method:	Least Squares	F-statistic:	966.3
Date:	Tue, 17 Dec 2019	Prob (F-statistic):	1.49e-158
Time:	08:47:02	Log-Likelihood:	-703.44
No. Observations:	233	AIC:	1419.
Df Residuals:	227	BIC:	1440.
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
chas	1.9041	1.303	1.461	0.145	-0.663	4.471

Model Attempt 5


```
In [77]: ▶ basetable6 = basetable5.drop(['chas'],axis=1)
basetable6.head()
```

Out[77]:

	rm	dis	ptratio	black	lstat	medv
0	6.575	4.0900	15.3	396.90	4.98	24.0
1	6.421	4.9671	17.8	396.90	9.14	21.6
2	6.998	6.0622	18.7	394.63	2.94	33.4
3	7.147	6.0622	18.7	396.90	5.33	36.2
4	6.012	5.5605	15.2	395.60	12.43	22.9


```
In [79]: ▶ sns.distplot(basetable5.rm, bins = bin_size, hist = True)
# Visualize the outliers in normal distribution
```

Out[79]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb777ca780>




```
In [80]: ▶ sns.distplot(basetable5.dis, bins = bin_size, hist = True)
# Visualize the outliers in normal distribution
```

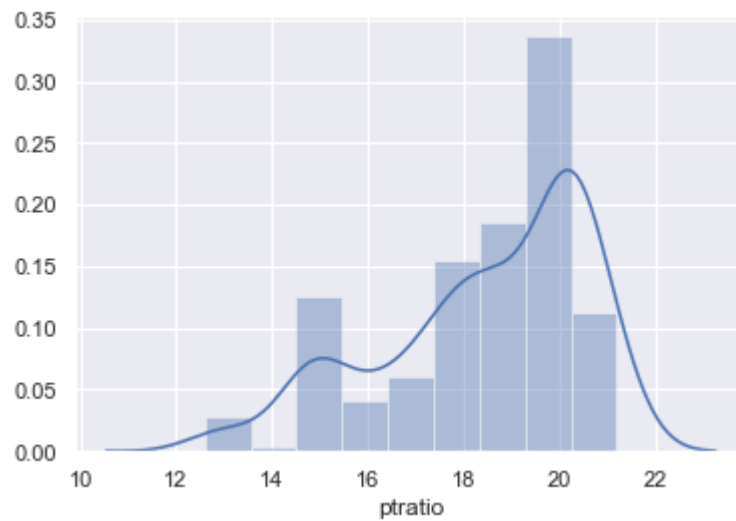
Out[80]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb77223cc0>



```
In [81]: ▶ #dis is negatively skewed
```

```
In [82]: ▶ sns.distplot(basetable5.ptratio, bins = bin_size, hist = True)
# Visualize the outliers in normal distribution
```


Out[82]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb77052240>



```
In [83]: ▶ #pt ratio is potitively skewed
```

```
In [84]: sns.distplot(basetable5.black, bins = bin_size, hist = True)
# Visualize the outliers in normal distribution
```


Out[84]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb77462588>



```
In [85]: #black is highly positively skewed
```

```
In [86]: sns.distplot(basetable5.lstat, bins = bin_size, hist = True)
# Visualize the outliers in normal distribution
```

Out[86]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb775a8f98>



In [87]: `#lstat is normal but negatively skewed`

In [88]: `#remove outliers`
`from scipy import stats`
`basetable6 = basetable5[(np.abs(stats.zscore(basetable5)) < 3).all(axis=1)]`
`print ("outliers removed beyond 3SD.")`

`basetable6.head()`


outliers removed beyond 3SD.

Out[88]:

	chas	rm	dis	ptratio	black	lstat	medv
0	0	6.575	4.0900	15.3	396.90	4.98	24.0
1	0	6.421	4.9671	17.8	396.90	9.14	21.6
2	0	6.998	6.0622	18.7	394.63	2.94	33.4
3	0	7.147	6.0622	18.7	396.90	5.33	36.2
4	0	6.012	5.5605	15.2	395.60	12.43	22.9


In [89]: `sns.distplot(basetable6.ptratio, bins = bin_size, hist = True)`
`# Visualize the outliers in normal distribution`

Out[89]: `<matplotlib.axes._subplots.AxesSubplot at 0x2bb77408048>`




```
In [90]: sns.distplot(basetable6.black, bins = bin_size, hist = True)
# Visualize the outliers in normal distribution
```

Out[90]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb7a849780>



```
In [96]: sns.distplot(basetable5.chas, bins = bin_size, hist = True)
# Visualize the outliers in normal distribution
```

Out[96]: <matplotlib.axes._subplots.AxesSubplot at 0x2bb7aaa4898>



```
In [97]: basetable6.head()
```

Out[97]:

	chas	rm	dis	ptratio	black	lstat	medv
0	0	6.575	4.0900	15.3	396.90	4.98	24.0
1	0	6.421	4.9671	17.8	396.90	9.14	21.6
2	0	6.998	6.0622	18.7	394.63	2.94	33.4
3	0	7.147	6.0622	18.7	396.90	5.33	36.2
4	0	6.012	5.5605	15.2	395.60	12.43	22.9

```
In [98]: basetable6 = basetable6.drop(['chas'],axis=1)
basetable6.head()
```

Out[98]:

	rm	dis	ptratio	black	lstat	medv
0	6.575	4.0900	15.3	396.90	4.98	24.0
1	6.421	4.9671	17.8	396.90	9.14	21.6
2	6.998	6.0622	18.7	394.63	2.94	33.4
3	7.147	6.0622	18.7	396.90	5.33	36.2
4	6.012	5.5605	15.2	395.60	12.43	22.9

```
In [99]: #get Predictor Dataframe
predictor = basetable6.drop('medv', axis = 1)
print(" Dependent variable : 'medv' Column removed from features")
predictor.head()

#get Target Dataframe
target = basetable6['medv']
target.head()

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(predictor, target, test_s

# Success
print("Training and testing split by 70/30 was successful")

print("Training Predictor dimension :",x_train.shape)
print("Training Target dimension :",y_train.shape)
print("Test Predictor dimension :",x_test.shape)
print("Test Target dimension :",y_test.shape)
```

```
Dependent variable : 'medv' Column removed from features
Training and testing split by 70/30 was successful
Training Predictor dimension : (200, 5)
Training Target dimension : (200,)
Test Predictor dimension : (87, 5)
Test Target dimension : (87,)
```

```
In [94]: #import library
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

#Model Training
lm5 = LinearRegression(fit_intercept=True,normalize=False)
print("Parameters of Linear Regressor function : ",lm5.get_params())

#Model Training
lm5.fit(x_train,y_train)


#Predict
y_pred = lm5.predict(x_test)
print("Total number of predicted values = ",y_pred.shape)

# The coefficients
plt.plot(lm5.coef_)
```

Parameters of Linear Regressor function : <bound method BaseEstimator.get_params of LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)>

Total number of predicted values = (87,)

Out[94]: [<matplotlib.lines.Line2D at 0x2bb7a9c6438>]



```
In [100]: ### Model Evalutation
```

```
In [101]: #Model Evaluation

from math import sqrt


#Calculate root-mean-square error (RMSE):
print("R-Squared for the above model : ",r2_score(y_test,y_pred)*100,"%")

#Calculate R-squared for the Model:
print("\\nroot-mean-square error (RMSE) for the model is : ",sqrt(mean_squared_error(y_test,y_pred))
```

R-Squared for the above model : -84.6901264244536 %

root-mean-square error (RMSE) for the model is : 11.520263468475049

Now, this is bad. Ha ha ha :)

In [104]:  basetable6.describe()

Out[104]:

	rm	dis	ptratio	black	lstat	medv
count	287.000000	287.000000	287.000000	287.000000	287.000000	287.000000
mean	6.273425	3.880460	18.418118	377.752753	11.941289	22.845645
std	0.626856	1.981831	2.132160	43.153169	6.422150	8.329401
min	4.368000	1.169100	12.600000	100.190000	1.730000	5.000000
25%	5.883500	2.262050	17.150000	379.540000	6.925000	18.100000
50%	6.208000	3.331700	18.700000	392.780000	10.420000	21.800000
75%	6.591500	5.344000	20.200000	396.900000	15.470000	25.000000
max	8.337000	9.222900	21.200000	396.900000	31.990000	50.000000

In []: 